



matloff sep. dir. for md's, and add R^2

History

1 contributor

826 lines (642 sloc) | 33.1 KB

...

The 10-Page Machine Learning Book

(The title here alludes to Andriy Burkov's excellent work, *The Hundred-Page Machine Learning Book*. Note too my own forthcoming book, *The Art of Machine Learning: Algorithms+Data+R*.)

Here we give an overview of the most widely used predictive methods in statistical/machine learning (ML). For each one, we present

- background
- overview of how it works
- function in the R package, **qeML** (readers without R background or who simply wish to acquire an overview of ML may skip the R code without loss of comprehension of the text)

Contents

- [Example](#)
- [Regression and classification problems, regression functions](#)
- [k-Nearest Neighbors](#)

- [K-NN edge bias](#)
- [Random forests](#)
 - [RF edge bias](#)
- [Boosting](#)
- [Linear model](#)
- [Logistic model](#)
- [Polynomial-linear models](#)
- [Shrinkage methods for linear/generalized linear models](#)
 - [LASSO for feature selection](#)
- [Support Vector Machines](#)
- [Neural networks](#)
- [Overfitting](#)
- [Which ML method to use?](#)

Notation

For convenience, we'll let Y denote the variable to be predicted, i.e. the response variable, and let X denote the set of predictor variables/features. (ML people tend to use the term *features*, while in stat, the term *predictors* is common. In applied fields, e.g. economics or psychology, some use the term *independent variables*.)

We develop our prediction rule from available *training data*, consisting of n data points, denoted by $(X_1, Y_1), \dots, (X_n, Y_n)$. We wish to predict new cases $(X_{\text{new}}, Y_{\text{new}})$ in the future, in which X_{new} is known but Y_{new} needs to be predicted.

So, we may wish to predict human weight Y from height X , or from height and age in a 2-component vector X . Say we have the latter situation, and data on $n = 100$ people. Then for instance X_{23} would be the vector of height and age for the 23rd person in our training data, and Y_{23} would be that person's weight.

The vector X may include *indicator* variables, which have values only 1 or 0. We may for instance predict weight from height, age and gender, the latter being 1 for female, 0 for male.

If Y represents a binary variable, we represent it as an indicator variable. In the famous Pima Diabetes dataset in the [UCI Machine Learning Repository](#), 1 means diabetic, 0 means not.

If Y itself is categorical, we represent it by several indicator variables, one for each category. In another disease-related UCI dataset, Y is status of a person's vertebrae condition; there are three classes: normal (NO), disk hernia (DH), or spondilolysthesis (SP). $Y = (1,0,0)$ for a normal person, for instance. Thus Y can be a vector too.

Running example

The package's built-in dataset **mlb** consists of data on major league baseball players (courtesy of the UCLA Dept. of Statistics).

Here is a glimpse of the data:

```
> data(mlb)
> head(mlb)
```

	Name	Team	Position	Height	Weight	Age	PosCategory
1	Adam_Donachie	BAL	Catcher	74	180	22.99	Catcher
2	Paul_Bako	BAL	Catcher	74	215	34.69	Catcher
3	Ramon_Hernandez	BAL	Catcher	72	210	30.78	Catcher
4	Kevin_Millar	BAL	First_Baseman	72	210	35.43	Infielder
5	Chris_Gomez	BAL	First_Baseman	73	188	35.71	Infielder
6	Brian_Roberts	BAL	Second_Baseman	69	176	29.39	Infielder

The R package's qe*-series functions

Here "qe" stands for "**quick and easy**," and we really mean that!

The functions have a simple, uniform interface, and most importantly, **require no setup**. To fit an SVM model, say, one simply calls **qeSVM**, no preparation calls to define the model etc.

The call form is

```
model fit <- qe<model name>(<data name>,<Y name>)
```

As noted, no prior calls are needed to define the model, etc.

Example

Let's predict weight from height and age, using two methods, k-Nearest Neighbor and random forests. (Details of the methods will be explained shortly; for now, let's just see how to invoke them.)

```
mlb1 <- mlb[,4:6] # columns for height, weight and age
knnout <- qeKNN(mlb1, 'Weight') # fit k-Nearest Neighbor model
rfout <- qeRF(mlb1, 'Weight') # fit random forests model
```

Most methods have *hyperparameters*, values that can be used to tweak the analysis in various ways. In **qeML**, default values of hyperparameters are set but can be overridden.

Prediction of new cases is equally easy, in the form

```
predict(<model fit>, <new X value>)
```

E.g. to predict the weight of a new player of height 70 and age 28, using our random forests model created above, run

```
> predict(rfout, c(70, 28))
      2
184.1626
```

Such a player would be predicted to weigh about 184 pounds.

The data is partitioned into a training set and a holdout set, with the model being fit on the former and then tested on the latter.

```
> rfout$testAcc # mean absolute prediction error
[1] 15.16911
```

So, on average, our predictions are off by about 15 pounds. What about the k-NN model?

```
> knnout$testAcc
[1] 13.20277
```

Seems to be better, though we should try other values of the hyperparameters.

(The size of the holdout set size can be set differently from the default, or suppressed entirely.)

Regression and classification problems, regression functions

Prediction applications in which Y is a continuous variable, say weight, or at least ordinal, are called *regression settings*. Applications in which Y is categorical, i.e. Y is a factor variable in R, say predicting the player's position (e.g. pitcher) are *classification settings*.

Somewhat confusingly, both settings make use of the *regression function*, $m(t) = E(Y | X = t)$, the mean value of Y in the subpopulation defined by $X = t$. If say we are predicting weight in the **mlb** data, then for instance $m(71,23)$ would be the mean weight among all players of height 71 inches and 23 years old. To predict the weight of a new player, say height 77 and age 19, we use $m(77,19)$.

In classification problems, Y is converted to a set of indicator variables. For the position 'pitcher' in the **mlb** data, we would have $Y = 1$ or 0 , depending on whether the player is a pitcher or not. (Position is in column 3 of the original dataset.) Then $E(Y | X = t)$ reduces to $P(Y = 1 | X = t)$, the probability that the player is a pitcher given the player's height, weight and age, say. We would typically find probabilities for each position, then guess the one with the highest probability.

In other words, the regression function $m(t)$ is central to both regression and classification settings. The statistical/machine learning methods presented here amount to ways to estimate $m(t)$. The methods are presented below in an order that shows connection between them.

Even though each ML method has its own special *tuning parameters* or *hyperparameters*, used to fine-tune performance, they all center around the regression function $m(t)$.

The **qe** series function sense whether the user is specifying a regression setting or a classification setting, by noting whether the Y variable (second argument) is numeric or an R factor.

ML predictive methods

We now present the "30,000 foot" view of the major statistical/machine learning methods.

k-Nearest Neighbors

This method was originally developed by statisticians, starting in the 1950s and 60s.

It's very intuitive. To predict, say, the weight of a new player of height 72 and age 25, we find the k closest players in our training data to (72,25), and average their weights. This is our estimate of $m(72,25)$, and we use it as our prediction. The default value of k in **qeKNN** is 25.

The **qeKNN** function wraps **kNN** in **regtools**. The main hyperparameter is the number of neighbors k . As with any hyperparameter, the user aims to set a "Goldilocks" level, not too big, not too small. Setting k too small will result in our taking the average of just a few Y values, too small a sample. On the other hand, too large a value for k will some distant data points may be used that are not representative.

To choose k , we could try various values, run **qeKNN** on each one, then use the value that produced the best (i.e. smallest) **testAcc**.

If Y is an indicator variable, its average works out to be the proportion of 1s. This will then be the estimated probability that $Y = 1$ for a new case. If that is greater than 0.5 in the neighborhood of the new case, we guess $Y = 1$ for the new case, otherwise 0.

If Y is categorical, i.e. an R factor as in the case of predicting player position in the **mlb** dataset, we then find a probability for each category in this manner, and guess Y to be whichever category has the highest probability.

The hyperparameter k as default value 25. Let's see if, say, 10 is better:

```
replicMeans(50,"qeKNN(mlb1,'Weight')$testAcc")
[1] 13.61954
attr(,"stderr")
[1] 0.1346821
replicMeans(50,"qeKNN(mlb1,'Weight',k=10)$testAcc")
[1] 14.25737
attr(,"stderr")
[1] 0.1298224
```

Since the holdout set is randomly generated, we did 50 runs in each case. The average test accuracy over 50 runs is printed out, along with a standard error for the figure. (1.96 times the standard error will be the radius of an approximate 95% confidence interval.) Changing k to 10 reduced accuracy.

The **qeML** function **qeFit** makes it easier to try various values of the hyperparameters.

K-NN edge bias

One potential problem is bias at the edge of a neighborhood. Say we are predict weight from height, and a new case involving a very tall person. Most data points in the neighborhood of this particular height value will be for people who are shorter than the new case. Those neighbors are thus likely to be lighter than the new case, making our predicted value for that case biased downward.

Rare for k-NN implementations, **qeKNN** has a remedy for this bias. Setting the argument **smoothingFtn = loclin** removes a linear trend within the neighborhood, and may improve predictive accuracy for new cases that are located near the edge of the training set.

Random forests

This method stems from *decision trees*, which were developed mainly by statisticians in the 1980s, and which were extended to random forests in 1990s. Note too the related (but unfortunately seldom recognized) *random subspaces* work of Tin Kam Ho, who did her PhD in computer science.

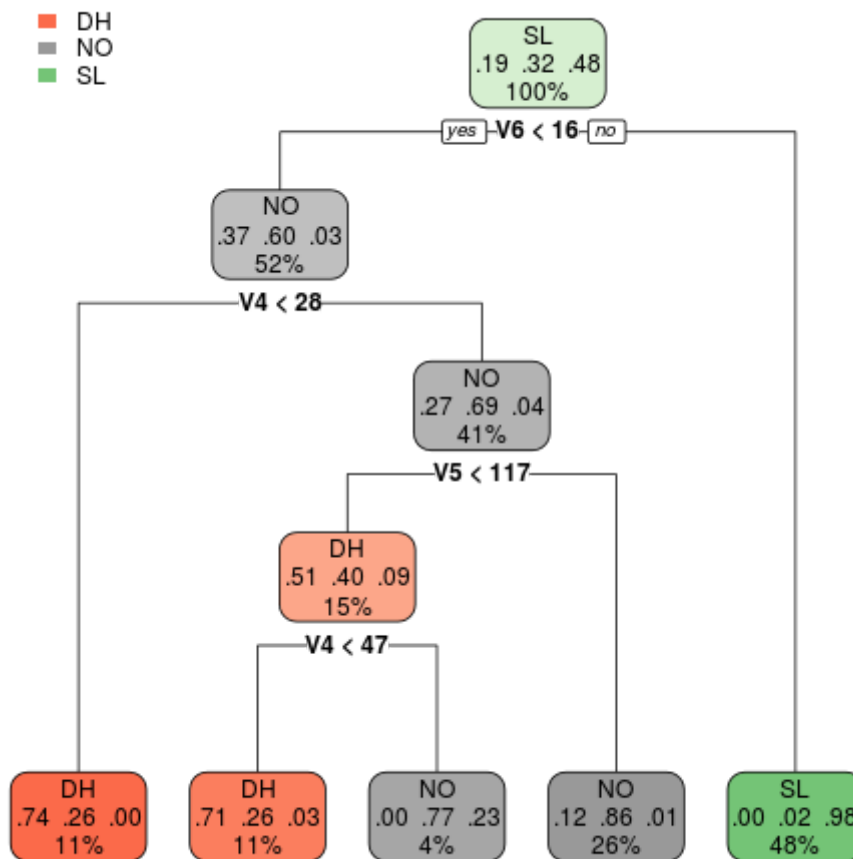
This is a natural extension of k-NN, in that it too creates a neighborhood and averages Y values within the neighborhood. However, it does so in a different way, creating tree structures.

Say in some dataset we are predicting blood pressure from height, weight and age. We first ask whether the height is above or below a certain threshold. After that, we ask whether weight is above or below a certain (different) threshold. This partitions height-weight space into 4 sectors. We then might subdivide each sector according to whether age is above or below a threshold, now creating 8 sectors of height-weight-age space. Each sector is now a "neighborhood." To predict a new case, we see which neighborhood it belongs to, then take our prediction to be the average Y value among training set points in that neighborhood.

The word *might* in the above paragraph alludes to the fact that the process may stop early, if the current subdivision is judged by the algorithm to be fine enough to produce good accuracy. So we might end up using only height and weight, not age.

And one generally wants to avoid having neighborhoods (*nodes* in the tree) that don't have many data points; this is controlled by a hyperparameter, say setting a minimum number of data points per node; if a split would violate that rule, then don't split. Of course, if we set out threshold too high, we won't do enough splits, so again we need to try to find a "Goldilocks" level.

Here is an example using the vertebrae dataset. There are 6 predictor variables, named V1 through V6, consisting of various bone measurements. A single tree is shown.



At the root, if the variable V6 in our new case to be predicted is < 16 , we go left, otherwise right. Say we go left. Then if $V4 < 28$ in the new case, we go left again, getting to a leaf, in which we guess DH. The 0.74 etc. mean that for the training data that happen to fall into that leaf, 74% of them are in class DH, 26% are NO and 0% are SL. So we guess DH, based on there being an estimated 74% chance that this new case is of type DH..

In RF, many trees are generated at random, using for instance random orders of entry of the features. To predict a new case, we find its predictions in the various trees, then average them. For indicator or categorical Y, each try gives a prediction, and we predict whatever class is most common among the trees.

Clearly, the order in which the predictor variables are evaluated (e.g. height, weight and age in the **mlb** data) can matter a lot. So, more than one tree is constructed, with random orders. The number of trees is another hyperparameter. Each tree gives us a prediction for the unknown Y . In a regression setting, those predictions are averaged to get our final prediction. In a classification setting, we see which class was predicted most often among all those trees.

The thresholds used at each node are determined through a complicated process, depending on which implementation of RF one uses.

The **qeRF** function wraps the function of the same name in the **randomForests** package.

RF edge bias

The package also offers other implementations of RF, notably **qeRFgrf**, as follows.

The leaves in any tree method, such as random forests and boosting, are essentially neighborhoods, different in structure from those of k -NN, but with similar properties. In particular, they have an "edge bias" problem similar to that described for k -NN above. In the case of random forests, the **qeRFgrf** function (which wraps the **grf** package) deals with the same bias problem via removal of a linear trend.

Boosting

This method has been developed both by CS and statistics people. The latter have been involved mainly in *gradient* boosting, the technique used here.

The basic idea is to iteratively build up a sequence of trees, each of which is an improved update of the last. At the end, all the trees are combined, with more weight given to the more recent trees.


The **qeGBoost** wraps **gbm** in the package of the same name. It is gradient tree-based, with hyperparameters similar to the random forests case, plus a *learning rate*. The latter controls the size of iteration steps; more on this below.

The package also offers other implementations of boosting.

Linear model

This of course is the classical linear regression model, invented 200 years ago (!) and developed by statisticians.

For motivation, below is a graph of mean weight vs. height for the **mlb** data. (There are many players for each height level, and we find and plot the mean weight at each height.)



The means seem to lie near a straight line. (Remember, though, that these are sample means.) That suggests modeling $m(t)$ is a linear function.

For example, a model for mean weight, given height and age, would be

$$m(\text{height}, \text{age}) = \beta_0 + \beta_1 \text{height} + \beta_2 \text{age}$$

for unknown population constants β_i , which are estimated from our training data, using the classic *least-squares* approach. Our estimates of the β_i , denoted b_i , are calculated by minimizing

$$\sum_i [\text{weight}_i - (b_0 + b_1 \text{height}_i + b_2 \text{age}_i)]^2$$

This is a simple calculus problem. We find the partial derivatives of the sum of squares with respect to the b_i , and set them to 0. This gives us 3 equations in 3 unknowns, and since these equations are linear, it is easy to solve them for the b_i .

There are no hyperparameters here.

This model is mainly for regression settings, though some analysts use it in classification. If used in conjunction with polynomials (see below), this may work as well or better than the logistic model (see below).

The function **qeLin** wraps the ordinary **lm**. It mainly just calls the latter, but does some little fixess.

Logistic model

This is a generalization of the linear model, developed by statisticians and economists.

This model is only for classification settings. As noted, since Y is now 1 or 0, its mean becomes the probability of 1. Since $m(t)$ is now a probability, we need it to have values in the interval $[0,1]$. This is achieved by feeding a linear model into the *logistic function*, $l(u) = (1 + \exp(-u))^{-1}$, which does take values in $(0,1)$.

So for instance, to predict whether a player is a catcher ($Y = 1$ if yes, $Y = 0$ if no), we fit the model

$$P(\text{catcher} \mid \text{height, weight, age}) = m(\text{height, weight, age}) = 1 / [1 + \exp\{-(\beta_0 + \beta_1 \text{height} + \beta_2 \text{weight} + \beta_3 \text{age})\}]$$

The β_i are estimated from the sample data, using a technique called *iteratively reweighted least squares*.

The function **qeLogit** wraps the ordinary R function **glm**, but adds an important feature: **glm** only handles the 2-class setting, e.g. catcher vs. non-catcher. The **qeLogit** handles the c -class situation via the *One vs. All* method (applicable to any ML algorithm):

it calls **glm** one class at a time, generating c **glm** outputs. When a new case is to be predicted, it is fed into each of the c **glm** outputs, yielding c probabilities. It then predicts the new case as whichever class has the highest probability.

Here is an example using the UCI vertebrae data;

```
> library(fdm2id)
> data(spine)
> str(spine)
'data.frame': 310 obs. of 8 variables:
 $ V1      : num 39.1 53.4 43.8 31.2 48.9 ...
 $ V2      : num 10.1 15.9 13.5 17.7 20 ...
 $ V3      : num 25 37.2 42.7 15.5 40.3 ...
 $ V4      : num 29 37.6 30.3 13.5 28.9 ...
 $ V5      : num 114 121 125 120 119 ...
 $ V6      : num 4.56 5.99 13.29 0.5 8.03 ...
```

```

$ Classif2: Factor w/ 2 levels "AB","NO": 1 1 1 1 1 1 1 1 1 1 ...
$ Classif3: Factor w/ 3 levels "DH","NO","SL": 1 1 1 1 1 1 1 1 1 1 ...
> spine <- spine[, -7] # skip 2-class example
> u <- qeLogit(spine, 'Classif3')
> u$testAcc
[1] 0.1935484
> u$baseAcc
[1] 0.5053763

> table(spine$Classif3)

DH NO SL
60 100 150

```

If we were to not use the body measurements V1 etc., we would always guess SL, resulting in an error rate of about 50%. By making use of the measurement data, we can reduce the misclassification rate to about 19%.

Let's try a prediction. Consider someone like the first patient in the dataset but with V6 = 6.22. What would our prediction be?

```

> newx <- spine[1, -7] # omit "Y"
> newx$V6 <- 6.22
> predict(u, newx)
$predClasses
[1] "DH"

$probs
          DH          NO          SL
[1,] 0.7432193 0.2420913 0.01468937

```

We would predict the DH class, as our estimated probability for the lass is 0.74, the largest among the three classes.

Some presentations describe the logistic model as "modeling the logarithm of the odds of $Y = 1$ vs. $Y = 0$ as linear." While this is correct, it is less informative, in our opinion. Why would we care about a logarithm being linear? The central issue is that the logistic function models a probability, just what we need.

Polynomial-linear models

Some people tend to shrink when they become older. Thus we may wish to model a tendency for people to gain weight in middle age but then lose weight as seniors, a nonlinear relation. We could try a quadratic model:

$$m(\text{height}, \text{age}) = \beta_0 + \beta_1 \text{height} + \beta_2 \text{age} + \beta_3 \text{age}^2$$

where presumably $\beta_3 < 0$.

We may even include a height X age product term, allowing for interactions. Polynomials of degree 3 and so on could also be considered. The choice of degree is a hyperparameter; in **qeSVM** it is named, of course, **degree**.

This would seem nonlinear, but that would be true only in the sense of being nonlinear in age. It is still linear in the β_i -- e.g. if we double each β_i in the above expression, the value of the expression is doubled -- so **qeLin** can be used, or **qeLogit** for classification settings.

Forming the polynomial terms by hand would be tedious, especially since we would also have to do this for predicting new cases. Instead, we use **qePolyLin** (regression setting) and **qePolyLog** (classification). They make use of the package **polyreg**.

Polynomial models can in many applications hold their own with the fancy ML methods. One must be careful, though, about overfitting, just as with any ML method. In particular, polynomial functions tend to grow rapidly near the edges of one's dataset, causing both bias and variance problems.

Shrinkage methods for linear/generalized linear models

Some deep mathematical theory implies that in linear models it may be advantageous to shrink the estimated b_i . *Ridge* regression and the LASSO do this in a mathematically rigorous manner. Each of them minimizes the usual sum of squared prediction errors, subject to a limit being placed on the size of the b vector; for ridge, the size is defined as the sum of the b_i^2 , while for the LASSO it's the sum of $|b_i|$.

The function **qeLASSO** wraps **cvglmnet** in the **glmnet** package. The main hyperparameter **alpha** specifies ridge (0) or the LASSO (1, the default).

There are various other shrinkage methods, such as Minimax Convex Penalty (MCP). This and some others are available via **qeNCVregCV**, which wraps the **ncvreg** package.

Shrinkage methods are often also applied to other ML algorithms.

The LASSO tends to produce solutions in which some of the b_i are 0. Thus it is popular as a tool for predictor variable selection.

LASSO for feature selection

The LASSO tends to produce solutions in which some of the b_i are 0. Thus it is popular as a tool for predictor variable selection.

Here is an example using **pef**, a dataset from the US Census, included with **qeML**. The features consist of age, education, occupation and gender. Those last three are categorical variables, which are converted to indicator variables. Let's predict wage income.

```
> w <- qeLASSO(pef, 'wageinc')
> w$coefs
11 x 1 sparse Matrix of class "dgCMatrix"
              s1
(Intercept) -8983.986
age          254.475
educ.14      9031.883
educ.16      9182.592
occ.100     -2707.388
occ.101     -1029.592
occ.102      3795.166
occ.106         .
occ.140         .
sex.1        4472.672
wkswrkd      1178.889
```

There are six occupations (this dataset is just for programmers and engineers),

```
> levels(pef$occ)
[1] "100" "101" "102" "106" "140" "141"
```


thus five indicator variables. The LASSO gave coefficients of 0 for occupations 106 and 140, so we might decide not to use them, even if we ultimately use some other ML algorithm.

Support Vector Machines

These were developed originally in the AI community, and later attracted interest among statisticians. They are used mainly in classification settings.

Say in the baseball data we are predicting catcher vs. non-catcher, based on height and weight. We might plot a scatter diagram, with height on the horizontal axis and weight on the vertical, using red dots for the catchers and blue dots for the non-catchers. We might draw a line that best separates the red and blue dots, then predict new cases by observing which side of the line they fall on. This is what SVM does (with more predictors, the line become a plane or hyperplane).

Here is an example, using the Iris dataset built in to R:



There are 3 classes, but we are just predicting setosa species (shown by + symbols) vs. non-setosa (shown by boxes) here. Below the solid line, we predict setosa, otherwise non-setosa.

SVM philosophy is that we'd like a wide buffer separating the classes, called the *margin*, denoted by the dashed lines. Data points lying on the edge of the margin are termed *support vectors*, so called because if any other data point were to change, the margin would not change.

In most cases, the two classes are not linearly separable. So we allow curved boundaries, implemented through polynomial (or similar) transformations to the data. The degree of the polynomial is a hyperparameter.

Another hyperparameter is **cost**: Here we allow some data points to be within the margin. The cost variable is roughly saying how many exceptions we are willing to accept.

The **qeSVM** function wraps **svm** in the **e1071** package. The package also offers various other implementations.

Neural networks

These were developed almost exclusively in the AI community.

An NN consists of *layers*, each of which consists of a number of *neurons* (also called *units* or *nodes*). Say for concreteness we have 10 neurons per layer. The output of the first layer will be 10 linear combinations of the predictor variables/features, essentially 10 linear regression models. Those will be fed into the second layer, yielding 10 "linear combinations of linear combinations," and so on.


In regression settings, the outputs of the last layer will be averaged together to produce our estimated $m(t)$. In the classification case with c classes, our final layer will have c outputs; whichever is largest will be our predicted class.

"Yes," you say, "but linear combinations of linear combinations are still linear combinations. We might as well just use linear regression in the first place." True, which is why there is more to the story: *activation functions*. Each output of a layer is fed into a function $A(t)$ before being passed on to the next layer; this is for the purpose of allowing nonlinear effects in our model of $m(t)$.

For instance, say we take $A(t) = t^2$ (not a common choice in practice, but a simple one to explain the issues). The output of the first layer will be quadratic functions of our features. Since we square again at the outputs of the second layer, the result will be 4th-degree polynomials in the features. Then 8th-degree polynomials come out of the third layer, and so on.

One common choice for $A(t)$ is the logistic function $l(u)$ we saw earlier. Another popular choice is ReLU, $r(t) = \max(0, t)$. No matter what we choose for $A(t)$, the point is that we have set up a nonlinear model for $m(t)$.

Here's an example, again with the Vertebrae data: The predictor variables V_1 , V_2 etc. for a new case to be predicted enter on the left, and the predictions come out the right; whichever of the 3 outputs is largest, that will be our predicted class.



The first layer consists of V1 through V6. The second layer, our only *hidden* layer here, has three neurons. Entering on the left of each neuron is a linear combination of V1 through V6. The outputs are fed into A(t) and then to the third layer.

Hyperparameters for NNs include the number of layers, the number of units per layer (which need not be the same in each layer), and the activation function. Most NN software also includes various other hyperparameters.

The linear combination coefficients, shown as numeric labels in the picture, are known as *weights*. How are they calculated? Again least squares is used, minimizing


$$\sum_i (Y_i - \text{finaloutput}_i)^2$$

Let n_w denote the number of weights. This can be quite large, even in the millions. Moreover, the n_w equations we get by setting the partial derivatives to 0 are not linear.

Thus this is no longer a "simple" calculus problem. Iterative methods must be used, and it can be extremely tricky to get them to converge. Here's why:

Though far more complex than in the linear case, we are still in the calculus realm. We compute the partial derivatives of the sum of squares with respect to the n_w weights, and set the results to 0s. So, we are finding roots of a very complicated function in n_w dimensions, and we need to do so iteratively.

A simplified version of the iteration process is as follows. Consider the function f graphed below:



There is an overall minimum at approximately $x = 2.2$. This is termed the *global minimum*. But there is also a *local minimum*, at about $x = 0.4$; that term means that this is the minimum value of the function only for points near---"local to"--- 0.4. Let's give the name x_0 to the value of x at the global minimum.

Denote our guess for x_0 at iteration i by g_i . Say our initial guess $g_0 = 1.1$.

The tangent line is pointing upward to the right, i.e. has positive slope, so it tells us that by going to the left we will go to smaller values of the function. We do want smaller values, but in this case, the tangent is misleading us. We should be going to the right, towards 2.2, where the global minimum is.

You can see that if our current guess were near 2.2, the tangent line would guide us in the right direction. But we see above that it can send us in the wrong direction. One remedy (among several typically used in concert) is to not move very far in the direction the tangent line sends us. The idea behind this is, if we are going to move in the wrong direction, let's limit our loss. The amount we move is called the *step size* in general math, but the preferred ML term is the *learning rate*. And, this is yet another hyperparameter.

So, NNs are arguably the most complex of all the methods described here, and tend to use huge amounts of computing time, even weeks!

The **qeNeural** function allows specifying the numbers of layers and neurons per layer, and the number of iterations. It wraps **krsFit** from the **regtools** package, which in turn wraps the R **keras** package (and there are further wraps involved after that).

Since any continuous activation function can be approximated by a polynomial, the outputs of the layers are approximately polynomials, of higher and higher degree with each layer. Thus they are subject to the same issues on the edges of the dataset as were described for polynomial models above.

Overfitting

Up to a point, the more complex a model is, the greater its predictive power. "More complex" can mean adding more predictor variables, using a higher-degree polynomial, adding more layers etc.

As we add more and more complexity, the *model bias* will decrease, meaning that our models become closer to the actual $m(t)$, in principle. But the problem is that at the same time, the *variance* of a predicted Y_{new} is increasing.

Say again we are predicting human weight from height, with polynomial models. With a linear model, we use our training data to estimate two coefficients, b_0 and b_1 . With a quadratic model, we estimate three, and estimate four in the case of a cubic model and so on. But it's same training data in each case, and intuitively we are "spreading the data thinner" with each more complex model. As a result, the standard error of our predicted Y_{new} (estimated standard deviation) increases.

Hence the famous Bias-Variance Tradeoff. If we use too complex a model, the increased variance overwhelms the reduction in bias, and our predictive ability suffers. We say we have *overfit*.

So there is indeed a "Goldilocks" level of complexity, an optimal polynomial degree, optimal number of nearest neighbors, optimal *network architecture* (configuration of layers and neurons), and so on. How do we find it?

Alas, **there are no magic answers here**. But various approaches have been developed that one can try. See the qeML vignettes **Feature_Selection** and **Overfitting**.

Which ML method to use?

The usual answer given for this question is, "The best method is very dependent on which dataset you have." True, but are there some general principles?

ML people often distinguish between "tabular" and "nontabular" data. A disease diagnosis application may be considered tabular--rows of patients, and columns of measurements such as blood glucose, body temperature and so on. On the other hand, a facial recognition application is considered nontabular, which is confusing, since it too has the form of a table--one row per picture, and one column per pixel position.

Where the two kinds of applications may differ, though, is whether there are general monotonic relationships between Y and the features. In predicting diabetes, say, the higher the glucose level, the more likely the patient is diabetic. Though monotonic relations may exist in image classification, they maybe more localized.

The point then, is that generally "monotonic" applications may be best served by ML methods that have either linear components or low-degree polynomials, such as linear or generalized linear models, SVM and polynomial versions of these. For other applications, one might try k-NN, NNs etc.